

Irfan Ahmed  
ICP 5  
Due Date: 9/28/2019

a. What you learned in the ICP.

I found this topic to be the most enjoyable so far. One of the main takes away from this ICP was how to logically setup my layers. After my data was train and tested on the dataset, it was normalized. Now the question arises, what should I perform next? I learned that applying one-hot encode would be the best option. One-hot encoding can be applied to the integer representation and doing so a new binary variable is added for each unique integer value. Also, few other concepts like convolution layer, pooling layer, and dropout method were experimented with to better understand the material. Different types of activation and optimizer were used to better understand what purpose they serve, and why one is used over the other. The importance of finding the right batch size and epochs can also affect the accuracy of the results.

b. ICP description what was the task you were performing

1. One of the first task was to upload the dataset and perform train/test on it.
2. Then the X and Y, train/test were normalized between 0 and 1 by dividing by 255.
3. One-hot encode was performed on the y\_train, Y\_test.
4. Next, the model was layed out using the Keras model, which acts as the blueprint.
5. The first layer of the model was the convolutional layer. This layer takes in the input and runs convolutional filters on them
6. Next, the drop layer was added to prevent overfitting, and batch normalization was also applied. Activation('relu') was used.
7. Polling layer was then added to make the classifier more robust.
8. Additional convolution layers were added to increase the number of filters.
9. After the convolution was finished the data was flattened and added dropout layer.
10. Denser layer was added specified the number of dense layers for each step.
11. In the final layer the number of classed was passed comparable to the number of neurons.
12. The model was finally compiled, and model fit was performed on the X\_train, and y\_train.
13. Then five new images were valeted to see the accuracy of the model.
14. In the final step 4 hyper-parameters were changed to see how it would affect the overall model accuracy(EX: Epochs, Batch Size, Dropout, and optimizers)

c. Challenges that you faced.

One of the difficult parts was understanding the different layers and how many to add. Adding can hurt the overall model and finding that sweet spot can be difficult. Mostly understanding the code and what is doing can be the most difficult part at times.

d. Screen shots that shows the successful execution of each required step of your code

- 1) Successfully executing the code and changing 4 hyperparameters in the model.
- 2) Validating the model on 5 new images (that are not present in the data set and are not used in training)
- 3) Providing the logical explanation of the changes that you made to hyper parameters and over all

Importing the required libaraies and packages.

```
[1] !pip install -q tf-nightly
import matplotlib.pyplot as plt
import numpy as np
import os
import PIL
import tensorflow as tf

from tensorflow import keras
from tensorflow.keras import layers
#from tensorflow.keras.models import Sequential
import numpy
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten, BatchNormalization, Activation
from keras.layers.convolutional import Conv2D, MaxPooling2D
from keras.constraints import maxnorm
from keras.utils import np_utils
from keras.datasets import cifar10
```

Setting random see for reproducibity

```
[2] seed = 25
```

```
[2] seed = 25
```

Loading the dataset. Which can be done by using load\_data()function

```
[3] (X_train, y_train), (X_test, y_test) = cifar10.load_data()
```

Looking at the shape of the data for all the previous varibales declared when the data was loaded

```
[4] print("The X_shape of the data: ", "size:", " M X M: ", "RBG Value:", '\n')

print("      x_train shape: ",X_train.shape, '\n')
print("      x_test shape: ",X_test.shape, '\n')

print("The X_shape of the data: ", "size:", "label:", '\n')

print("      x_train shape: ",y_train.shape, '\n')
print("      x_test shape: ",y_test.shape, '\n')
```

↳ The X\_shape of the data: size: M X M: RBG Value:

```
      x_train shape: (50000, 32, 32, 3)
      x_test shape: (10000, 32, 32, 3)
```

The X\_shape of the data: size: label:

```
      x_train shape: (50000, 1)
      x_test shape: (10000, 1)
```

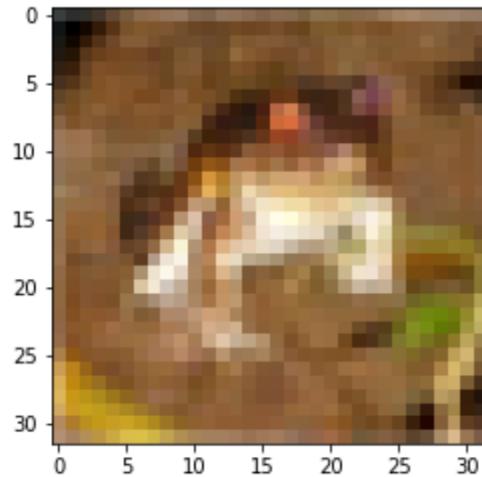
printing each label. It represents each class.

+ Code

```
▶ print(y_test)  
↳ [[3]  
 [8]  
 [8]  
 ...  
 [5]  
 [1]  
 [7]]
```

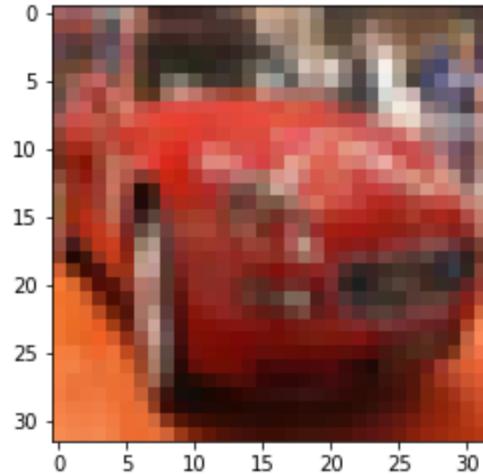
looking at the image at a random location

```
[6] plt.imshow(X_train[0])  
↳ <matplotlib.image.AxesImage at 0x7f849b4d9e48>
```



```
[7] plt.imshow(X_train[5])
```

```
▶ plt.imshow(X_train[5])  
⇨ <matplotlib.image.AxesImage at 0x7f849bf892e8>
```



```
[8] #num_classes = 10  
#cifar10_classes = ["airplane", "automobile", "bird", "cat", "deer",  
 "#dog", "frog", "horse", "ship", "truck"]
```

Normalizing the data.

The input values are between 0 and 255.

```
[9] X_train = X_train.astype('float32')  
X_test = X_test.astype('float32')  
X_train /= 255.0 # dividing by 255  
X_test /= 255.0
```

```
▶ X_train = X_train.astype('float32')
X_test = X_test.astype('float32')
X_train /= 255.0 # dividing by 255
X_test /= 255.0
```

Converts a class vector (integers) to binary class matrix. This will be done for y\_train and y\_test. Re

```
[10] y_train = np_utils.to_categorical(y_train,10)
y_test = np_utils.to_categorical(y_test,10)
```

```
[11] print("y_train",'\n')
print(y_train[0:5], '\n')
print("-----")
print("y_test", '\n')
print(y_test[0:5], '\n')
```

⇨ y\_train

```
[[0. 0. 0. 0. 0. 0. 1. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 0. 0. 0. 1.]
 [0. 0. 0. 0. 0. 0. 0. 0. 0. 1.]
 [0. 0. 0. 0. 1. 0. 0. 0. 0. 0.]
 [0. 1. 0. 0. 0. 0. 0. 0. 0. 0.]]
```

-----

y\_test

```
[[0. 0. 0. 1. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 0. 0. 1. 0.]
 [0. 0. 0. 0. 0. 0. 0. 0. 1. 0.]
 [1. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 1. 0. 0. 0.]]
```

specifying the number of classes present in the dataset. The purpose of doing this allows us know in the final layer

```
▶ class_num = y_test.shape[1]
```

Building the model(Blue print)

The convolutional layer simply put takes in a image which is the input and assigns the importance the image. This allows the it to differentiate them.

```
[13] model = Sequential()
      model.add(Conv2D(32, (3,3), input_shape=X_train.shape[1:], padding='same'))
```

The rectified linear activation function or ReLU for short is a piecewise linear function that will output it will output zero. It has become the default activation function for many types of neural networks train and often achieves better performance

```
[14] model.add(Activation('relu'))
```

Next, the dropout layer is added. The reason for adding this is to prevent neural networks from overfitting. The dropout layer is to reduce the capacity or thinning the network during training. Another thing to note maybe needed when it is utilized.

Purpose of adding BatchNormalization is to standardize the inputs, meaning it will have mean of zero and standard deviation of one.

```
[15] model.add(Dropout(0.2)) #adding 20% dropout
      model.add(BatchNormalization()) #used to standardize the input
      model.add(Conv2D(64, (3, 3), padding = 'same'))#2D convolution layer
      model.add(Activation('relu'))
```

The pooling layer is added after the convolutional layer (ReLU). It is just a new layer that is added.

```
[16] model.add(MaxPooling2D(pool_size=(2,2))) #takes in 2x2 pooling window  
model.add(Dropout(0.2))  
model.add(BatchNormalization())
```

Adding another pooling layer. total of 2 more layers added and reason for this is because the imag

```
[17] model.add(Conv2D(64, (3, 3), padding='same'))  
model.add(Activation('relu'))  
model.add(MaxPooling2D(pool_size=(2, 2)))  
model.add(Dropout(0.2))  
model.add(BatchNormalization())
```

```
[18] model.add(Conv2D(128, (3, 3), padding='same'))  
model.add(Activation('relu'))  
model.add(Dropout(0.2))  
model.add(BatchNormalization())
```

Flatteing the layer This will convert the pooled features map into a single column which is passed to the final Dense layer.

```
[19] model.add(Flatten())  
model.add(Dropout(0.2))
```

The Dense layer is imported. Few things need to be addressed before doing so. First, the number of neurons in the dense layer decreases and will approach the same number of neurons as there are classes.

```
[20] model.add(Dense(256, kernel_constraint=maxnorm(3)))
    model.add(Activation('relu'))
    model.add(Dropout(0.2))
    model.add(BatchNormalization())
```

Repeating the code now decreasing the units from 256 to 128. Factor of 2

```
[21] model.add(Dense(256, kernel_constraint=maxnorm(3)))
    model.add(Activation('relu'))
    model.add(Dropout(0.2))
    model.add(BatchNormalization())
```

In the final stages of the layer the number of classes are passed for the number of neurons. class\_num is passed on line 12 of code. For the activation 'softmax' is used because it will select the neuron that have the highest probability.

```
[22] model.add(Dense(class_num))
    model.add(Activation('softmax'))
```

The number of epochs are trained and the optimizer that will be used. Epoch is an dataset that is passed through the neural networks at once. Epoch is generally the what we want to train the CNN

```
[23]
    epochs = 2
```

Finally, compiling the data and see the results.

```
[24] model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```
[25] print(model.summary())
```

↳ Model: "sequential"

Layer (type)	Output Shape	Param #
<hr/>		
conv2d (Conv2D)	(None, 32, 32, 32)	896
activation (Activation)	(None, 32, 32, 32)	0
dropout (Dropout)	(None, 32, 32, 32)	0
batch_normalization (BatchNo	(None, 32, 32, 32)	128
conv2d_1 (Conv2D)	(None, 32, 32, 64)	18496
activation_1 (Activation)	(None, 32, 32, 64)	0
max_pooling2d (MaxPooling2D)	(None, 16, 16, 64)	0
dropout_1 (Dropout)	(None, 16, 16, 64)	0
batch_normalization_1 (Batch	(None, 16, 16, 64)	256
conv2d_2 (Conv2D)	(None, 16, 16, 64)	36928

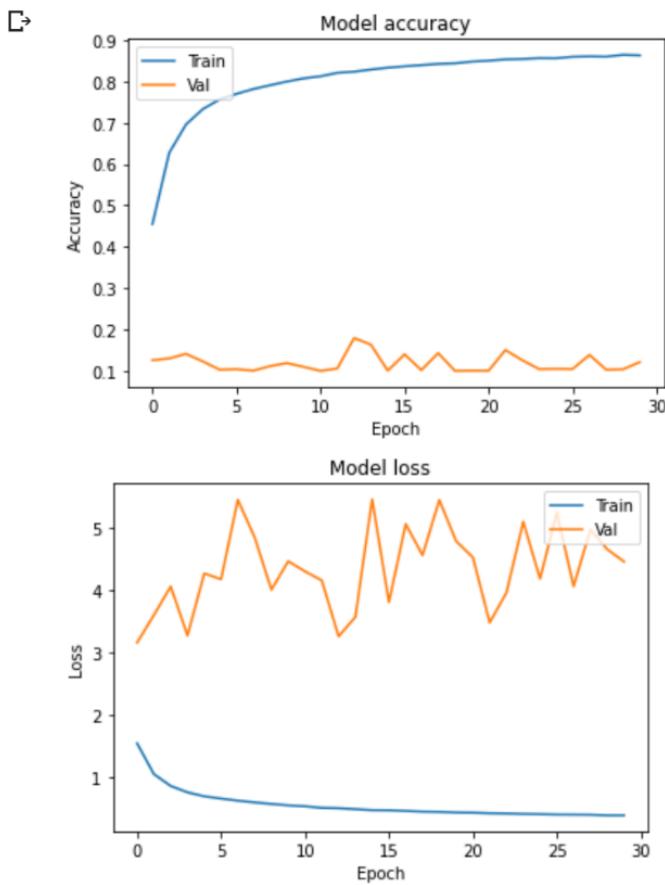


```
numpy.random.seed(seed)
data_history11 = model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=22)

↳ 782/782 [=====] - 412s 526ms/step - loss: 1.1162 - accuracy: Epoch 3/30
782/782 [=====] - 408s 522ms/step - loss: 0.8809 - accuracy: Epoch 4/30
782/782 [=====] - 410s 524ms/step - loss: 0.7557 - accuracy: Epoch 5/30
782/782 [=====] - 407s 521ms/step - loss: 0.6870 - accuracy: Epoch 6/30
782/782 [=====] - 407s 521ms/step - loss: 0.6424 - accuracy: Epoch 7/30
782/782 [=====] - 407s 520ms/step - loss: 0.6154 - accuracy: Epoch 8/30
782/782 [=====] - 408s 521ms/step - loss: 0.5753 - accuracy: Epoch 9/30
782/782 [=====] - 409s 523ms/step - loss: 0.5551 - accuracy: Epoch 10/30
782/782 [=====] - 413s 528ms/step - loss: 0.5210 - accuracy: Epoch 11/30
782/782 [=====] - 410s 525ms/step - loss: 0.5104 - accuracy: Epoch 12/30
782/782 [=====] - 405s 518ms/step - loss: 0.4790 - accuracy: Epoch 13/30
782/782 [=====] - 404s 517ms/step - loss: 0.4736 - accuracy: Epoch 14/30
782/782 [=====] - 405s 518ms/step - loss: 0.4537 - accuracy: Epoch 15/30
782/782 [=====] - 402s 514ms/step - loss: 0.4500 - accuracy: Epoch 16/30
782/782 [=====] - 404s 516ms/step - loss: 0.4397 - accuracy: Epoch 17/30
782/782 [=====] - 408s 522ms/step - loss: 0.4267 - accuracy: Epoch 18/30
782/782 [=====] - 406s 520ms/step - loss: 0.4209 - accuracy: Epoch 19/30
782/782 [=====] - 406s 519ms/step - loss: 0.4182 - accuracy: Epoch 20/30
782/782 [=====] - 406s 519ms/step - loss: 0.4030 - accuracy: Epoch 21/30
782/782 [=====] - 411s 526ms/step - loss: 0.4096 - accuracy: Epoch 22/30
```

```
[1]: plt.plot(data_history11.history['accuracy'])
plt.plot(data_history11.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Val'], loc='upper left')
plt.show()

plt.plot(data_history11.history['loss'])
plt.plot(data_history11.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Val'], loc='upper right')
plt.show()
```



```
[30] scores11 = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores11[1]*100))
```

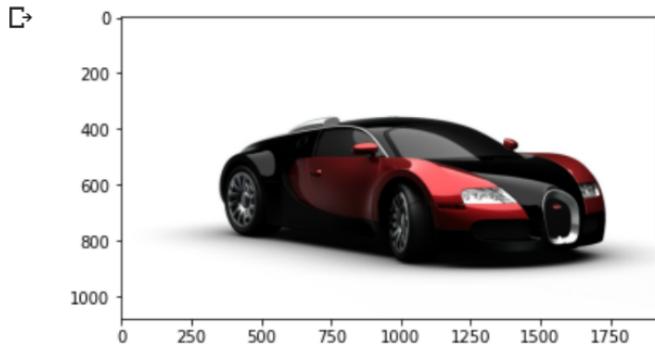
↳ Accuracy: 12.09%

Now validating the results with random images

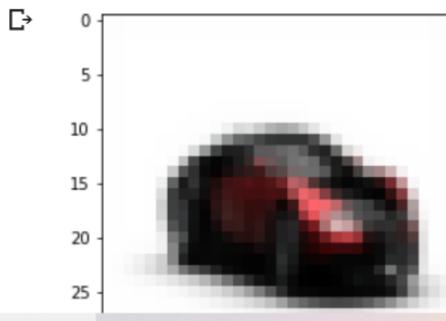
```
[31] #Load the data
from google.colab import files # Use to load data on Google Colab
uploaded = files.upload() # Use to load data on Google Colab
new_image = plt.imread("/content/car-49278_1920.jpg") #Read in the image
```

↳ Choose Files car-49278\_1920.jpg  
• car-49278\_1920.jpg(image/jpeg) - 106000 bytes, last modified: 9/26/2020 - 100% done  
Saving car-49278\_1920.jpg to car-49278\_1920.jpg

```
[32] img = plt.imshow(new_image)
```



```
[33] from skimage.transform import resize
resized_image = resize(new_image, (32,32,3))
img = plt.imshow(resized_image)
```



```
[34] predictions = model.predict(np.array( [resized_image] ))  
  
[35] predictions  
↳ array([[1.0230566e-02, 7.7509576e-01, 4.0212185e-03, 1.7838222e-01,  
         1.8251414e-04, 7.1143741e-03, 7.5949458e-03, 2.6186512e-04,  
         1.2788793e-02, 4.3276763e-03]], dtype=float32)  
  
[36] classification = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']  
  
[37] list_index = [0,1,2,3,4,5,6,7,8,9]  
    x = predictions  
    for i in range(10):  
        for j in range(10):  
            if x[0][list_index[i]] > x[0][list_index[j]]:  
                temp = list_index[i]  
                list_index[i] = list_index[j]  
                list_index[j] = temp  
    #Show the sorted labels in order from highest probability to lowest  
    print(list_index)  
  
↳ [1, 3, 8, 0, 6, 5, 9, 2, 7, 4]  
  
[38] i=0  
    for i in range(5):  
        print(classification[list_index[i]], ':', round(predictions[0][list_index[i]] * 100, 2), '%')  
  
↳ automobile : 77.51 %  
    cat : 17.84 %  
    ship : 1.28 %  
    airplane : 1.02 %  
    frog : 0.76 %
```

image 2

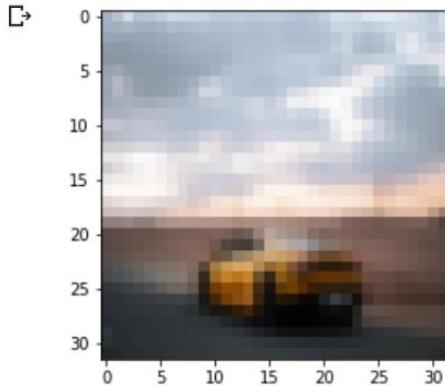
```
[39] #Load the data
from google.colab import files # Use to load data on Google Colab
uploaded = files.upload() # Use to load data on Google Colab
new_image = plt.imread("/content/car-5548242_1920.jpg") #Read in the image
```

↳ Choose Files car-5548242\_1920.jpg  
• car-5548242\_1920.jpg(image/jpeg) - 290156 bytes, last modified: 9/26/2020 - 100% done  
Saving car-5548242\_1920.jpg to car-5548242\_1920.jpg

```
[40] img = plt.imshow(new_image)
```



```
[41] from skimage.transform import resize
resized_image = resize(new_image, (32,32,3))
img = plt.imshow(resized_image)
```



```

[42] predictions = model.predict(np.array( [resized_image] ))
predictions

⇒ array([[2.1332918e-01, 2.7268991e-02, 6.1268936e-04, 9.6222095e-05,
       1.3771025e-04, 6.0589850e-06, 4.2500244e-05, 1.5197129e-05,
       6.9925004e-01, 5.9241354e-02]], dtype=float32)

+ Code

[43] list_index = [0,1,2,3,4,5,6,7,8,9]
x = predictions
for i in range(10):
    for j in range(10):
        if x[0][list_index[i]] > x[0][list_index[j]]:
            temp = list_index[i]
            list_index[i] = list_index[j]
            list_index[j] = temp
#Show the sorted labels in order from highest probability to lowest
print(list_index)

⇒ [8, 0, 9, 1, 2, 4, 3, 6, 7, 5]

[44] i=0
for i in range(5):
    print(classification[list_index[i]], ':', round(predictions[0][list_index[i]] * 100, 2), '%')

⇒ ship : 69.93 %
airplane : 21.33 %
truck : 5.92 %
automobile : 2.73 %
bird : 0.06 %

```

image 3

```

[45] #Load the data
from google.colab import files # Use to load data on Google Colab
uploaded = files.upload() # Use to load data on Google Colab
new_image = plt.imread("/content/cuba-1197800_1920.jpg") #Read in the image

img = plt.imshow(new_image)

```

image 3

```
[45] #Load the data
from google.colab import files # Use to load data on Google Colab
uploaded = files.upload() # Use to load data on Google Colab
new_image = plt.imread("/content/cuba-1197800_1920.jpg") #Read in the image

img = plt.imshow(new_image)
```

Choose Files cuba-1197800\_1920.jpg  
• cuba-1197800\_1920.jpg(image/jpeg) - 837906 bytes, last modified: 9/26/2020 - 100% done  
Saving cuba-1197800\_1920.jpg to cuba-1197800\_1920.jpg



```
[46] from skimage.transform import resize
resized_image = resize(new_image, (32,32,3))
img = plt.imshow(resized_image)

predictions = model.predict(np.array( [resized_image] ))
predictions
```

array([[1.3416114e-04, 1.3246661e-02, 1.0425618e-02, 1.5270984e-03,
 9.4200277e-01, 1.5187241e-03, 8.1295198e-05, 3.5432293e-03,
 2.6010284e-03, 2.4919322e-02]], dtype=float32)



```

▶ list_index = [0,1,2,3,4,5,6,7,8,9]
x = predictions
for i in range(10):
    for j in range(10):
        if x[0][list_index[i]] > x[0][list_index[j]]:
            temp = list_index[i]
            list_index[i] = list_index[j]
            list_index[j] = temp
#Show the sorted labels in order from highest probability to lowest
print(list_index)

i=0
for i in range(5):
    print(classification[list_index[i]], ':', round(predictions[0][list_index[i]] * 100, 2), '%')

⇨ [4, 9, 1, 2, 7, 8, 3, 5, 0, 6]
deer : 94.2 %
truck : 2.49 %
automobile : 1.32 %
bird : 1.04 %
horse : 0.35 %

```

image 4

```

[48] #Load the data
from google.colab import files # Use to load data on Google Colab
uploaded = files.upload() # Use to load data on Google Colab
new_image = plt.imread("/content/fiat-4322521_1920.jpg") #Read in the image

img = plt.imshow(new_image)

```

⇨ Choose Files fiat-4322521\_1920.jpg  
• fiat-4322521\_1920.jpg(image/jpeg) - 696404 bytes, last modified: 9/26/2020 - 100% done  
Saving fiat-4322521\_1920.jpg to fiat-4322521\_1920.jpg



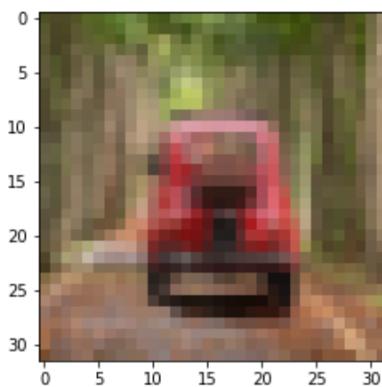
```

▶ from skimage.transform import resize
resized_image = resize(new_image, (32,32,3))
img = plt.imshow(resized_image)

predictions = model.predict(np.array( [resized_image] ))
predictions

```

⇒ array([[1.0125967e-04, 6.4570587e-03, 5.8513653e-04, 1.8047021e-03, 1.0990113e-03, 2.4159544e-04, 7.3753879e-05, 3.6307792e-05, 4.9243611e-03, 9.8467690e-01]], dtype=float32)



[50] list\_index = [0,1,2,3,4,5,6,7,8,9]

```

x = predictions
for i in range(10):
    for j in range(10):
        if x[0][list_index[i]] > x[0][list_index[j]]:
            temp = list_index[i]
            list_index[i] = list_index[j]
            list_index[j] = temp
#Show the sorted labels in order from highest probability to lowest
print(list_index)

```

i=0

```

for i in range(5):
    print(classification[list_index[i]], ':', round(predictions[0][list_index[i]] * 100, 2), '%')

```

⇒ [9, 1, 8, 3, 4, 2, 5, 0, 6, 7]

```

truck : 98.47 %
automobile : 0.65 %
ship : 0.49 %
cat : 0.18 %
deer : 0.11 %

```

image 5

```
#Load the data
from google.colab import files # Use to load data on Google Colab
uploaded = files.upload() # Use to load data on Google Colab
new_image = plt.imread("/content/ford-498244_1920.jpg") #Read in the image

img = plt.imshow(new_image)

Choose Files ford-498244_1920.jpg
• ford-498244_1920.jpg(image/jpeg) - 582919 bytes, last modified: 9/26/2020 - 100% done
Saving ford-498244_1920.jpg to ford-498244_1920.jpg
```

```
[52]
from skimage.transform import resize
resized_image = resize(new_image, (32,32,3))
img = plt.imshow(resized_image)

predictions = model.predict(np.array( [resized_image] ))
predictions
```

```
array([[1.2390652e-03, 1.7010628e-01, 6.8462832e-04, 8.4482197e-04,
       4.1197411e-05, 7.4195523e-05, 7.6745840e-04, 8.3668010e-06,
       6.3244629e-01, 1.9378777e-01]], dtype=float32)
```



```

list_index = [0,1,2,3,4,5,6,7,8,9]
x = predictions
for i in range(10):
    for j in range(10):
        if x[0][list_index[i]] > x[0][list_index[j]]:
            temp = list_index[i]
            list_index[i] = list_index[j]
            list_index[j] = temp
#Show the sorted labels in order from highest probability to lowest
print(list_index)

i=0
for i in range(5):
    print(classification[list_index[i]], ':', round(predictions[0][list_index[i]] * 100, 2), '%')

⇒ [8, 9, 1, 0, 3, 6, 2, 5, 4, 7]
ship : 63.24 %
truck : 19.38 %
automobile : 17.01 %
airplane : 0.12 %
cat : 0.08 %

```

4 hyper-parameters that were changed

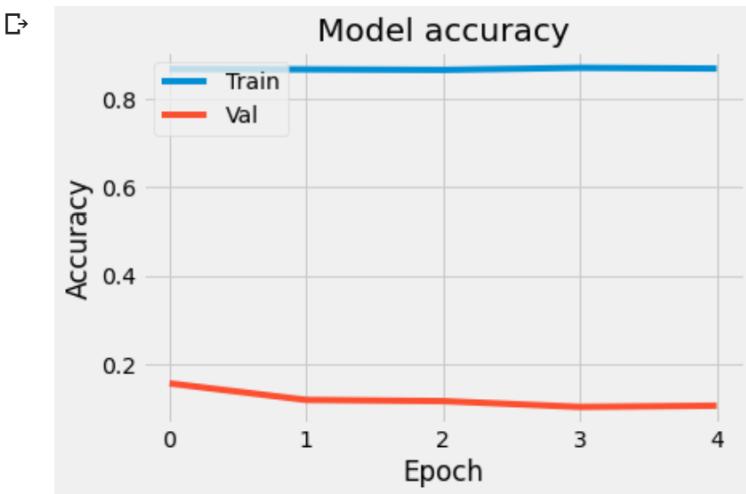
```

numpy.random.seed(seed)
data_history00 = model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=5, batch_size=64)

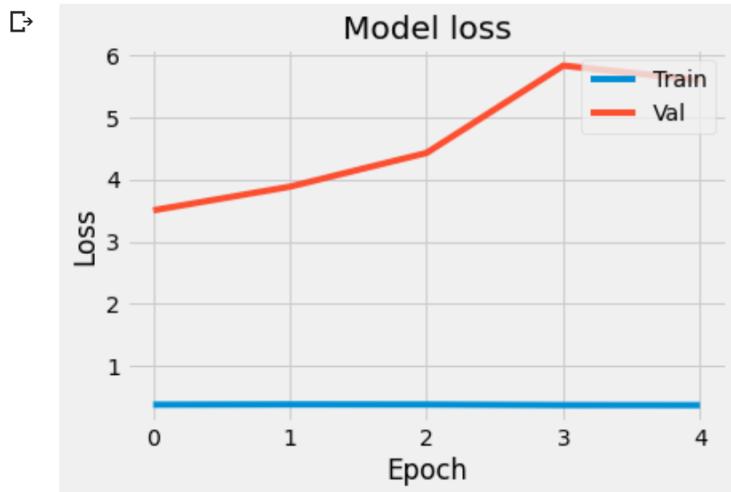
⇒ Epoch 1/5
782/782 [=====] - 407s 520ms/step - loss: 0.3779 - accuracy: 0.8672 - val_loss: 3.50
Epoch 2/5
782/782 [=====] - 408s 522ms/step - loss: 0.3819 - accuracy: 0.8664 - val_loss: 3.88
Epoch 3/5
782/782 [=====] - 410s 525ms/step - loss: 0.3806 - accuracy: 0.8653 - val_loss: 4.42
Epoch 4/5
782/782 [=====] - 409s 523ms/step - loss: 0.3705 - accuracy: 0.8700 - val_loss: 5.83
Epoch 5/5
782/782 [=====] - 407s 520ms/step - loss: 0.3691 - accuracy: 0.8683 - val_loss: 5.59

```

```
▶ import matplotlib.pyplot as plt  
plt.style.use('fivethirtyeight')  
  
[ ] plt.plot(data_history00.history['accuracy'])  
plt.plot(data_history00.history['val_accuracy'])  
plt.title('Model accuracy')  
plt.ylabel('Accuracy')  
plt.xlabel('Epoch')  
plt.legend(['Train', 'Val'], loc='upper left')  
plt.show()
```



```
▶ plt.plot(data_history00.history['loss'])
plt.plot(data_history00.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Val'], loc='upper right')
plt.show()
```



```
[ ] # Model evaluation
scores2 = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores2[1]*100))
```

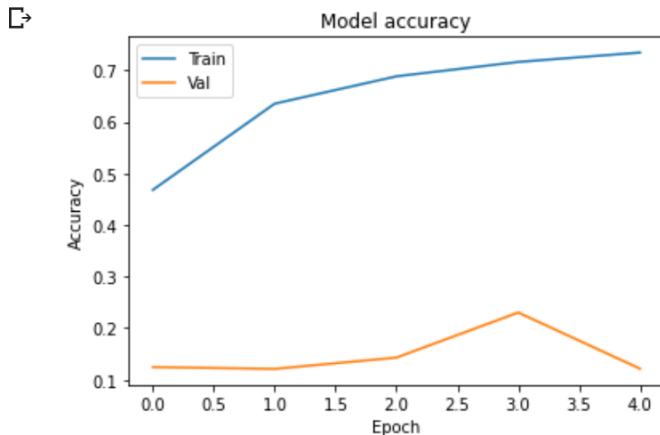
□ Accuracy: 10.56%

changing batch size to 32

```
▶ numpy.random.seed(seed)
data_history1 = model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=5, batch_size=32)

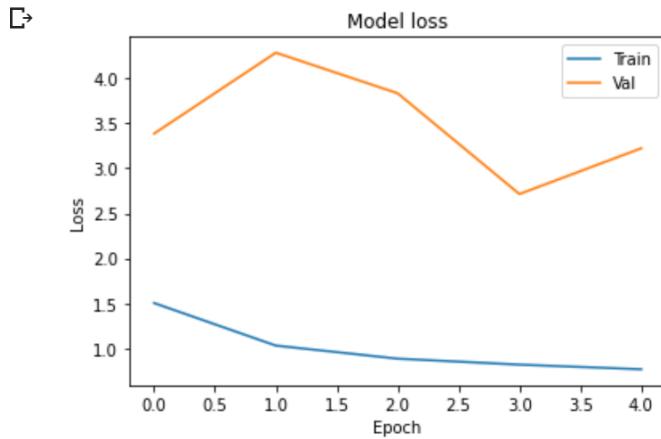
▷ Epoch 1/5
1563/1563 [=====] - 436s 278ms/step - loss: 1.8076 - accuracy: 0.3782 - val_loss: 3.0000 - val_accuracy: 0.1250
Epoch 2/5
1563/1563 [=====] - 433s 277ms/step - loss: 1.0773 - accuracy: 0.6168 - val_loss: 4.0000 - val_accuracy: 0.1250
Epoch 3/5
1563/1563 [=====] - 433s 277ms/step - loss: 0.8875 - accuracy: 0.6876 - val_loss: 3.0000 - val_accuracy: 0.1250
Epoch 4/5
1563/1563 [=====] - 433s 277ms/step - loss: 0.8112 - accuracy: 0.7154 - val_loss: 2.0000 - val_accuracy: 0.1250
Epoch 5/5
1563/1563 [=====] - 434s 277ms/step - loss: 0.7467 - accuracy: 0.7392 - val_loss: 3.0000 - val_accuracy: 0.1250
```

```
[41] plt.plot(data_history1.history['accuracy'])
plt.plot(data_history1.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Val'], loc='upper left')
plt.show()
```



Epoch

```
[42] plt.plot(data_history1.history['loss'])
    plt.plot(data_history1.history['val_loss'])
    plt.title('Model loss')
    plt.ylabel('Loss')
    plt.xlabel('Epoch')
    plt.legend(['Train', 'Val'], loc='upper right')
    plt.show()
```



```
[43] # Model evaluation
    scores3 = model.evaluate(X_test, y_test, verbose=0)
    print("Accuracy: %.2f%%" % (scores3[1]*100))
```

Accuracy: 12.15%

changing epochs to 6

```
▶ numpy.random.seed(seed)
    data_history3 = model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=6, batch_size=16)
```

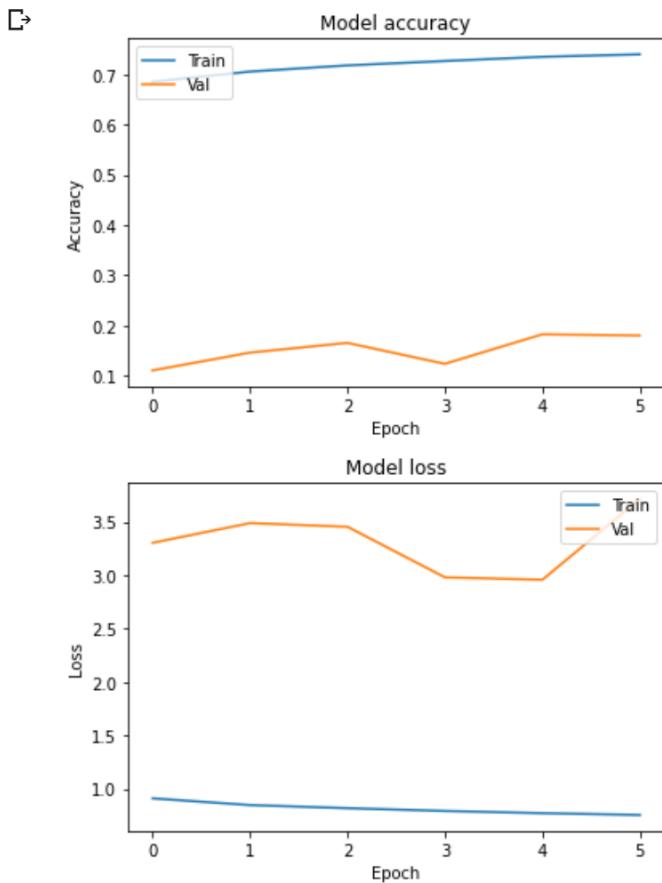
Epoch 1/6  
3125/3125 [=====] - 486s 155ms/step - loss: 0.9100 - accuracy: 0.6856 - val\_loss:  
Epoch 2/6  
3125/3125 [=====] - 483s 154ms/step - loss: 0.8469 - accuracy: 0.7059 - val\_loss:  
Epoch 3/6  
3125/3125 [=====] - 490s 157ms/step - loss: 0.8178 - accuracy: 0.7183 - val\_loss:  
Epoch 4/6  
3125/3125 [=====] - 488s 156ms/step - loss: 0.7920 - accuracy: 0.7270 - val\_loss:  
Epoch 5/6  
3125/3125 [=====] - 489s 157ms/step - loss: 0.7709 - accuracy: 0.7354 - val\_loss:  
Epoch 6/6  
3125/3125 [=====] - 488s 156ms/step - loss: 0.7541 - accuracy: 0.7403 - val\_loss:

```

▶ plt.plot(data_history3.history['accuracy'])
plt.plot(data_history3.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Val'], loc='upper left')
plt.show()

plt.plot(data_history3.history['loss'])
plt.plot(data_history3.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Val'], loc='upper right')
plt.show()

```



```
[46] # Model evaluation
    scores4 = model.evaluate(X_test, y_test, verbose=0)
    print("Accuracy: %.2f%%" % (scores4[1]*100))
```

```
↳ Accuracy: 17.97%
```

chaning optimizer using Adamax instead of Adam. Also, the learing rate is changed It is a variant of Adam based on the infin parameters follow those provided in the paper. Adamax is sometimes superior to adam, specially in models with embedding

```
[61] model.compile(loss='categorical_crossentropy', optimizer='adamax', metrics=['accuracy'])
    numpy.random.seed(seed)
    data_history6 = model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=5, batch_size=64)
```

```
↳ Epoch 1/5
782/782 [=====] - 399s 508ms/step - loss: 0.4548 - accuracy: 0.8434 - val_loss: 3.19
Epoch 2/5
782/782 [=====] - 398s 509ms/step - loss: 0.3904 - accuracy: 0.8643 - val_loss: 3.37
Epoch 3/5
782/782 [=====] - 397s 508ms/step - loss: 0.3584 - accuracy: 0.8737 - val_loss: 3.21
Epoch 4/5
782/782 [=====] - 397s 508ms/step - loss: 0.3360 - accuracy: 0.8821 - val_loss: 3.01
Epoch 5/5
782/782 [=====] - 397s 507ms/step - loss: 0.3171 - accuracy: 0.8889 - val_loss: 3.13
```

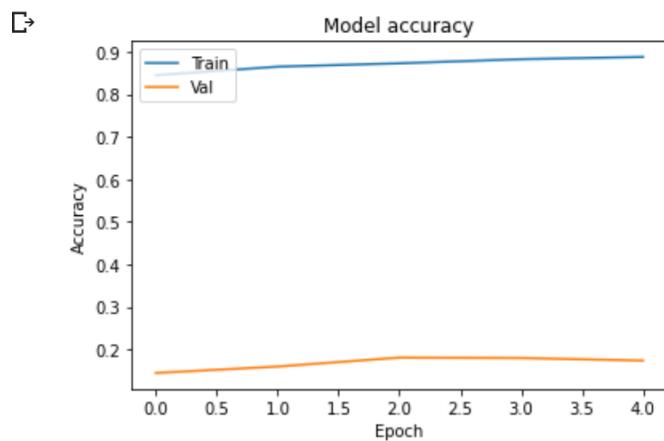
```
[48] numpy.random.seed(seed)
    data_history6 = model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=5, batch_size=64)
    #numpy.random.seed(seed)
    #data_history5 = model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=5, batch_size=64)
```

```
↳ Epoch 1/5
782/782 [=====] - 401s 511ms/step - loss: 0.7548 - val_loss: 2.7824
Epoch 2/5
782/782 [=====] - 400s 512ms/step - loss: 0.6258 - val_loss: 3.6033
Epoch 3/5
782/782 [=====] - 401s 512ms/step - loss: 0.6160 - val_loss: 3.8410
Epoch 4/5
782/782 [=====] - 401s 512ms/step - loss: 0.5665 - val_loss: 3.5670
Epoch 5/5
782/782 [=====] - 401s 513ms/step - loss: 0.5655 - val_loss: 2.9442
```

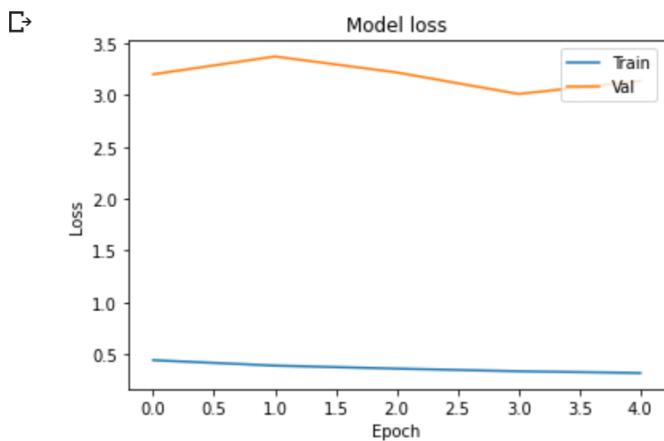
```
[62] history_dict = data_history6.history
     print(history_dict.keys())

    □ dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])

    ▶▶ plt.plot(data_history6.history['accuracy'])
        plt.plot(data_history6.history['val_accuracy'])
        plt.title('Model accuracy')
        plt.ylabel('Accuracy')
        plt.xlabel('Epoch')
        plt.legend(['Train', 'Val'], loc='upper left')
        plt.show()
```



```
plt.plot(data_history6.history['loss'])
plt.plot(data_history6.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Val'], loc='upper right')
plt.show()
```



```
[65] # Model evaluation
scores5 = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores5[1]*100))
```

Accuracy: 17.36%

chaning dropout to 0.3

```
model.add(Dense(256, kernel_constraint=maxnorm(3)))
model.add(Activation('relu'))
model.add(Dropout(0.3))
model.add(BatchNormalization())
```

```
[67] model.add(Dense(class_num))
model.add(Activation('softmax'))
```

```
[68] model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```
print(model.summary())

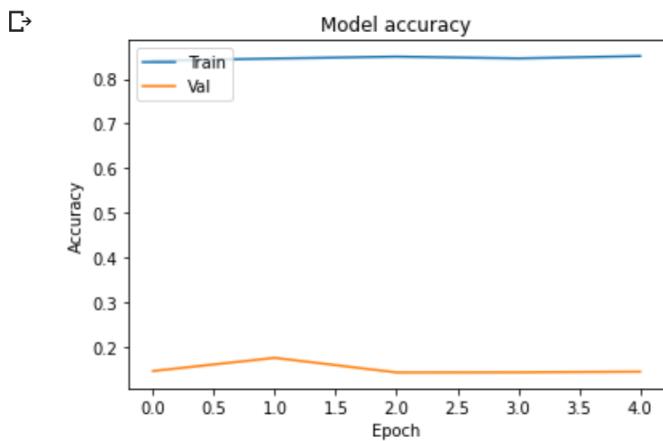
```

Layer (type)	Output Shape	Param #
conv2d_4 (Conv2D)	(None, 32, 32, 32)	896
activation_7 (Activation)	(None, 32, 32, 32)	0
dropout_7 (Dropout)	(None, 32, 32, 32)	0
batch_normalization_6 (Batch Normalization)	(None, 32, 32, 32)	128
conv2d_5 (Conv2D)	(None, 32, 32, 64)	18496
activation_8 (Activation)	(None, 32, 32, 64)	0
max_pooling2d_2 (MaxPooling2D)	(None, 16, 16, 64)	0
dropout_8 (Dropout)	(None, 16, 16, 64)	0
batch_normalization_7 (Batch Normalization)	(None, 16, 16, 64)	256
conv2d_6 (Conv2D)	(None, 16, 16, 64)	36928
activation_9 (Activation)	(None, 16, 16, 64)	0
max_pooling2d_3 (MaxPooling2D)	(None, 8, 8, 64)	0
dropout_9 (Dropout)	(None, 8, 8, 64)	0
batch_normalization_8 (Batch Normalization)	(None, 8, 8, 64)	256
conv2d_7 (Conv2D)	(None, 8, 8, 128)	73856
activation_10 (Activation)	(None, 8, 8, 128)	0
dropout_10 (Dropout)	(None, 8, 8, 128)	0
batch_normalization_9 (Batch Normalization)	(None, 8, 8, 128)	512
flatten_1 (Flatten)	(None, 8192)	0
dropout_11 (Dropout)	(None, 8192)	0
dense_3 (Dense)	(None, 256)	2097408
activation_11 (Activation)	(None, 256)	0
dropout_12 (Dropout)	(None, 256)	0
batch_normalization_10 (Batch Normalization)	(None, 256)	1024
dense_4 (Dense)	(None, 256)	65792

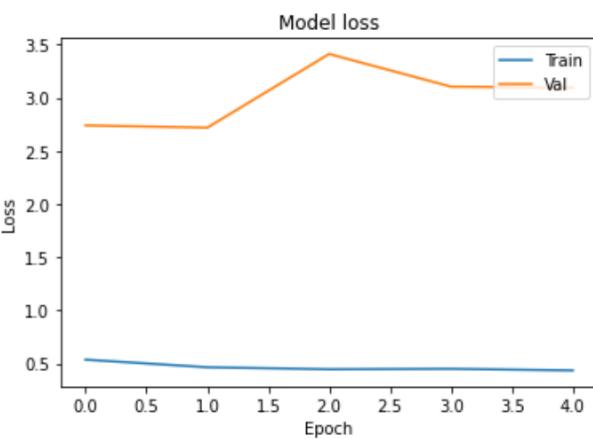
```
▶ numpy.random.seed(seed)
data_history7 = model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=5, batch_size=64)

▷ Epoch 1/5
782/782 [=====] - 402s 512ms/step - loss: 0.6218 - accuracy: 0.8215 - val_loss: 2.7
Epoch 2/5
782/782 [=====] - 400s 512ms/step - loss: 0.4504 - accuracy: 0.8529 - val_loss: 2.7
Epoch 3/5
782/782 [=====] - 398s 508ms/step - loss: 0.4251 - accuracy: 0.8589 - val_loss: 3.4
Epoch 4/5
782/782 [=====] - 399s 510ms/step - loss: 0.4265 - accuracy: 0.8559 - val_loss: 3.1
Epoch 5/5
782/782 [=====] - 399s 511ms/step - loss: 0.4128 - accuracy: 0.8570 - val_loss: 3.0
```

```
[71] plt.plot(data_history7.history['accuracy'])
plt.plot(data_history7.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Val'], loc='upper left')
plt.show()
```



```
[1]: plt.plot(data_history7.history['loss'])
plt.plot(data_history7.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Val'], loc='upper right')
plt.show()
```



```
[2]: # Model evaluation
scores6 = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores6[1]*100))
```

[2]: Accuracy: 14.39%

e. Output file link if applicable.

<https://github.com/UMKC-APL-BigDataAnalytics/icp5-irfancheemaa>

f. Video link (YouTube or any other publicly available video platform).

You should have access to the link. If there are any issues, please let me know.

<https://umkc.hosted.panopto.com/Panopto/Pages/Viewer.aspx?id=a7676aed-4d2a-4e69-95a5-ac440187923e>

g. Any inside about the data or the ICP in general

I enjoyed this topic and always wanted to learn more about it and how to implement it. It would be nice to learn few more variations of this method with smaller dataset.