

Irfan Ahmed

ICP 9

Due Date: 10/26/2020

Data Explanation:

For this report Mean Squared Error, Y intercept and slope were calculated based on the insurance dataset used. I was trying to look at the cost of insurance charges and compare it with sex, age, bmi, region, smoker, and children. There is some indication that smokers do tend to pay higher cost but that could also be attributed to the person age as well and where they live. The MSE for this observation was 0.501, Y intercept 0.00566 and the slopes for each of the variable were the following: [0.29254799 - 0.0116126 ,0.15938599, 0.03498638, 0.80153266, -0.0290586]. The r2_score was 0.754, which can be interpreted as less difference between the observed data and the fitted values. The MSE was 0.501, which could be the model needs improvement to get better understanding of if any further relationships can be assumed with the features and charges.

a. What you learned in the ICP

For this ICP the goal was to create a linear regression model. Linear regression can be used to find the relationship between continuous variables. Where one is the predictor/independent variable and the other is the dependent variable.

b. ICP description what was the task you were performing

1. Imported the needed packages/libraries.
2. Reading the csv file for insurance cost.
3. Looked at the data to see what could be removed or changed.
4. Checked for null values.
5. Converted string into numerical representation using LabelEncoder. The sex, region and smoker strings were changed to their appropriate presentation.
6. Used PCA on the dataset.
7. Created feature list and X and y variable for train and test purposes.
8. Both X and y datasets were standardized.
9. Training and testing method were applied.
10. Used LinearRegression() for analysis. The model was fitted for the X_train and y_train. Also, prediction value was calculated.
11. Finally, the intercept, coefficient and MSE were calculated.

c. Challenges that you faced

My data size did not match at first but after playing around with different variable I was able to fix it to where I could run my model. Since I was comparing multiple columns, I had difficulty trying to visual that in my code.

d. Screen shots that shows the successful execution of each required step of your code

Create a linear regression model in python using any dataset of your choice. For this model you can also create your own data. Find the best fit line in the data and calculate SSE (sum of square error) or MSE (Mean square error), Y intercept, and Slope for the relationship in data. Explain your findings and understanding of these terms in detail in the report.

Successfully executing the code with linear regression model and calculating following: a. SSE or MSE b. Y intercept c. Slope

Mean Absolute Error: 0.34193698828461927

Mean Squared Error: 0.25108181197832785

Root Mean Squared Error: 0.5010806441864701

Y Intercept: 0.005664780213938073

Slope: [0.29254799 -0.0116126 0.15938599 0.03498638 0.80153266 -0.0290586]

Importing the required packages/libraries

```
[1] #import packages/libraries
import pandas as pd
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
import numpy as np
from sklearn.linear_model import LinearRegression
```

Reading the insurance data

```
[2] data = pd.read_csv('/content/insurance.csv') #reading the csv file
```

```
[3] data.head(3) #looking at the first 5 data points
   age sex bmi children smoker region charges
0 19  female 27.90      0    yes southwest 16884.9240
1 18  male 33.77      1     no southeast 1725.5523
2 28  male 33.00      3     no southeast 4449.4620
```

standardize age, bmi, children

pca

```
[4] data.describe()
   age      bmi   children   charges
count 1338.000000 1338.000000 1338.000000
mean 39.207025 30.663397 1.094918 13270.422265
std 14.049980 6.098187 1.205493 12110.011237
min 18.000000 15.980000 0.000000 1121.873900
25% 27.000000 26.296250 0.000000 4740.287150
50% 39.000000 30.400000 1.000000 9382.033000
75% 51.000000 34.693750 2.000000 16639.912515
max 64.000000 53.130000 5.000000 63770.428010
```

```
[5] data.describe() []
   age      bmi   children   charges
count 1338.000000 1338.000000 1338.000000
mean 39.207025 30.663397 1.094918 13270.422265
std 14.049980 6.098187 1.205493 12110.011237
min 18.000000 15.980000 0.000000 1121.873900
25% 27.000000 26.296250 0.000000 4740.287150
50% 39.000000 30.400000 1.000000 9382.033000
75% 51.000000 34.693750 2.000000 16639.912515
max 64.000000 53.130000 5.000000 63770.428010
```

Checking for null values

```
[5] data.isnull().sum() #checking for null values
   age      0
   sex      0
   bmi      0
   children 0
   smoker   0
   region   0
   charges   0
dtype: int64
```

Will need to convert to number format for sex, region and smoker. 0 is smoker and 1 is not a smoker. 0 is female and 1 is male. value for region is based on the location.

```
[6] from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
le.fit(data.sex.drop_duplicates())
data.sex = le.transform(data.sex) #sex male or female
#print(data.sex)

le.fit(data.smoker.drop_duplicates())
data.smoker = le.transform(data.smoker) #smoker yes or no
#print(data.smoker)

le.fit(data.region.drop_duplicates()) #region
data.region = le.transform(data.region)
#print(data.region)
```

```

[7] data.head(5)

   age sex bmi children smoker region charges
0 19 0 27.900 0 1 3 16884.92400
1 18 1 33.770 1 0 2 1725.55230
2 28 1 33.000 3 0 2 4449.46200
3 33 1 22.705 0 0 1 21984.47061
4 32 1 28.880 0 0 1 3866.80520

[8] from sklearn.decomposition import PCA
pca = PCA(whiten=True)
pca.fit(data)
variance = pd.DataFrame(pca.explained_variance_ratio_)
np.cumsum(pca.explained_variance_ratio_)

array([0.099999851, 0.09999974, 0.09999998, 0.09999999, 1.        ,
       1.        ]))

Creating a features and X and y variables

[9] feature_cols = ['age', 'sex', 'bmi', 'children', 'smoker', 'region']

X = data[feature_cols]
y = data.charges

[11] #standardized X and y
X = preprocessing.scale(X)
y = preprocessing.scale(y)

Training the X and y datasets

[12] X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 1)

Using the linear regression model for analysis

[13] linear_model = LinearRegression()

[14] model_fit = linear_model.fit(X_train,y_train) #fitting the training and test data

[15] linear_model.score(X_train,y_train) #looking at the score
0.754083642384213

[16] pred = linear_model.predict(X_test) #prediction value

[17] from sklearn.metrics import r2_score
r2_score(y_test, pred) #R squared
0.740367716897532

For retrieving the slope (coefficient of x)

[18] linear_model.coef_ #gives you an array of weights estimated by linear regression. It is of shape (ntargets, nfeatures)
array([ 0.20254709, -0.0116126,  0.15938599,  0.03498638,  0.80153266,
       -0.0290586 ])

To retrieve the intercept

[19] linear_model.intercept_
0.005647480213938073

[20] from sklearn.metrics import mean_squared_error
from sklearn import metrics

[22] from sklearn import metrics
print('Mean Absolute Error: ', metrics.mean_absolute_error(y_test, pred))
print('Mean Squared Error: ', metrics.mean_squared_error(y_test, pred))
print('Root Mean Squared Error: ', np.sqrt(metrics.mean_squared_error(y_test, pred)))
print('Y intercept: ', linear_model.intercept_)
print('Slope: ', linear_model.coef_)

Mean Absolute Error: 0.3419369883801927
Mean Squared Error: 0.25100181377832705
Root Mean Squared Error: 0.501000641804701
Y Intercept: 0.005647480213938073
Slope: [ 0.29254799,  0.0116126,  0.15938599,  0.03498638,  0.80153266,  0.0290586 ]

```

e. Output file link if applicable

<https://github.com/UMKC-APL-BigDataAnalytics/icp9-irfancheemaa>

f. Video link (YouTube or any other publicly available video platform)

<https://umkc.hosted.panopto.com/Panopto/Pages/Viewer.aspx?id=9b7fe325-a1f8-4dcf-bf4a-ac6001505c68>

g. Any inside about the data or the ICP in general

I really enjoyed this ICP because it allowed me to create my own model. Learning how to properly setup each step to make sure the next line of code be appropriate to write and use.